

ASSIGNMENT-17.3

Name:Ch.Rohith reddy

Rollno:2303A52140

Batch:41

Task 1 Employee Dataset Cleaning

Task:

Use AI assistance to clean an employee information dataset.

Instructions:

- Handle missing values in salary, designation, and experience
- Remove duplicate employee IDs
- Standardize text values (job titles in lowercase)
- Encode categorical features (designation, employment_type)

Prompt :

Clean an employee dataset.

Instructions:

- Handle missing values in salary (median), designation (mode), and experience (mean)
- Remove duplicate records based on employee_id
- Convert designation text to lowercase
- Encode categorical columns: designation and employment_type
- Output the cleaned dataset with no missing values and encoded features

Code :

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder

# ---- Sample Dataset ----
np.random.seed(42)
df = pd.DataFrame({
    "employee_id": np.random.randint(100, 200, 50),
    "name": [f"Emp{i}" for i in range(50)],
    "salary": np.random.choice([30000, 40000, 50000, 60000, np.nan], 50),
    "designation": np.random.choice(["Manager", "Developer", "Analyst",
None], 50),
    "experience": np.random.choice([1, 2, 3, 4, 5, np.nan], 50),
```

```
    "employment_type": np.random.choice(["Full-time", "Part-time",  
"Contract"], 50)  
})
```

```

# ---- 1. Handle Missing Values ----
df['salary'] = df['salary'].fillna(df['salary'].median())
df['designation'] = df['designation'].fillna(df['designation'].mode()[0])
df['experience'] = df['experience'].fillna(df['experience'].mean())

# ---- 2. Remove Duplicate Employee IDs ----
df = df.drop_duplicates(subset='employee_id')

# ---- 3. Standardize Text Values ----
df['designation'] = df['designation'].str.lower()

# ---- 4. Encode Categorical Features ----
le1 = LabelEncoder()
le2 = LabelEncoder()

df['designation'] = le1.fit_transform(df['designation'])
df['employment_type'] = le2.fit_transform(df['employment_type'])

# ---- Output ----
print("Cleaned Dataset:")
print(df.head())

print("\nShape:", df.shape)
print("\nMissing Values:\n", df.isnull().sum())

```

Output:

```

... Cleaned Dataset:
   employee_id  name  salary  designation  experience  employment_type
0          151  Emp0  40000.0             0           2.0                0
1          192  Emp1  60000.0             0           2.0                1
2          114  Emp2  30000.0             0           4.0                2
3          171  Emp3  60000.0             0           5.0                1
4          160  Emp4  40000.0             0           3.0                1

Shape: (37, 6)

Missing Values:
employee_id      0
name             0
salary           0
designation       0
experience        0
employment_type  0
dtype: int64

```

Justification:

- Missing Values Handling
 - *Salary* → *Median*: Salary distributions are often skewed; median is robust to outliers.
 - *Designation* → *Mode*: Most frequent job title maintains categorical consistency.
 - *Experience* → *Mean*: Numerical field; mean preserves overall distribution.
- Duplicate Removal
 - Removing duplicate `employee_id` ensures data integrity and avoids biased analytics.
- Text Standardization
 - Converting designations to lowercase eliminates case-sensitive inconsistencies (e.g., “Manager” vs “manager”).
- Categorical Encoding
 - Machine learning models require numerical inputs; encoding converts text into usable numeric form.

Task 2 – Student Performance Data Preprocessing

Task:

Preprocess student performance data using AI assistance.

Instructions:

- Convert exam_date column into datetime format
- Extract new features (month, semester, day)
- Normalize total_marks
- Detect and handle outliers in marks using IQR or Z-score

Prompt:

Preprocess a student performance dataset.

Instructions:

- Convert exam_date to datetime format
- Extract new features: month, day, semester (1 for Jan–Jun, 2 for Jul–Dec)
- Normalize total_marks using MinMax scaling
- Detect and remove outliers using IQR method
- Output the transformed dataset with new features and normalized marks

Code:

```

import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

# ---- Sample Dataset ----
np.random.seed(0)
dates = pd.date_range(start="2024-01-01", periods=60)

df =
    pd.DataFrame({ "student_id":
        range(1, 61),
        "exam_date": np.random.choice(dates, 60),
        "total_marks": np.random.randint(200, 1000, 60)
    })

# Add some outliers
df.loc[5, 'total_marks'] = 1500
df.loc[10, 'total_marks'] = 50

# ---- 1. Convert exam_date to datetime ----
df['exam_date'] = pd.to_datetime(df['exam_date'])

# ---- 2. Extract Features ----
df['month'] = df['exam_date'].dt.month
df['day'] = df['exam_date'].dt.day
df['semester'] = df['month'].apply(lambda x: 1 if x <= 6 else 2)

# ---- 3. Normalize total_marks ----
scaler = MinMaxScaler()
df['total_marks'] = scaler.fit_transform(df[['total_marks']])

# ---- 4. Detect & Remove Outliers (IQR) ----
Q1 = df['total_marks'].quantile(0.25)
Q3 = df['total_marks'].quantile(0.75)
IQR = Q3 - Q1

df = df[(df['total_marks'] >= Q1 - 1.5 * IQR) &
        (df['total_marks'] <= Q3 + 1.5 * IQR)]

# ---- Output ----
print("Processed Dataset:")

```

```
print(df.head())
```

```
print("\nShape:", df.shape)
```

```
print("\nColumns:", df.columns)
```

Output:

```
... Processed Dataset:
   student_id exam_date total_marks month day semester
0           1 2024-02-14    0.229655     2  14         1
1           2 2024-02-17    0.122759     2  17         1
2           3 2024-02-23    0.191724     2  23         1
3           4 2024-01-01    0.191724     1   1         1
4           5 2024-01-04    0.140000     1   4         1

Shape: (59, 6)

Columns: Index(['student_id', 'exam_date', 'total_marks', 'month', 'day', 'semester'], dtype='object')
```

Justification:

- Datetime Conversion
 - Enables time-based feature extraction and chronological analysis.
- Feature Extraction (month, day, semester)
 - Captures academic patterns and seasonality (e.g., exam trends across semesters).
- Normalization (MinMax Scaling)
 - Scales marks to a uniform range (0-1), preventing feature dominance in ML models.
- Outlier Removal (IQR)
 - Eliminates extreme values that can distort model training and statistical analysis.

Task 3 – Student Health Data Preparation

Task:

Clean and preprocess student health records using AI scripts.

Instructions:

- Handle missing values in height, weight, BMI
- Encode categorical features (medical_condition, fitness_status)
- Scale numerical features using standard scaling
- Split dataset into training and testing sets

Prompt:

Prepare a student health dataset for machine learning.

Instructions:

- Handle missing values in height, weight, and BMI using mean
- Encode categorical features: medical_condition and fitness_status
- Apply standard scaling to numerical features (height, weight, BMI)
- Split dataset into training and testing sets (80-20)
- Output processed dataset ready for modeling

Code:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

# ---- Sample Dataset ----
np.random.seed(1)

df =
    pd.DataFrame({ "student_id":
        range(1, 51),
        "height": np.random.choice([150, 160, 170, 180, np.nan], 50),
        "weight": np.random.choice([50, 60, 70, 80, np.nan], 50),
        "BMI": np.random.choice([18, 22, 25, 30, np.nan], 50),
        "medical_condition": np.random.choice(["None", "Asthma", "Diabetes"],
50),
        "fitness_status": np.random.choice(["Fit", "Unfit"], 50)
    })

# ---- 1. Handle Missing Values ----
df['height'] = df['height'].fillna(df['height'].mean())
df['weight'] = df['weight'].fillna(df['weight'].mean())
df['BMI'] = df['BMI'].fillna(df['BMI'].mean())

# ---- 2. Encode Categorical Features ----
le1 = LabelEncoder()
le2 = LabelEncoder()

df['medical_condition'] = le1.fit_transform(df['medical_condition'])
df['fitness_status'] = le2.fit_transform(df['fitness_status'])
```

```

# ---- 3. Scale Numerical Features ----
scaler = StandardScaler()
df[['height', 'weight', 'BMI']] = scaler.fit_transform(df[['height',
'weight', 'BMI']])

# ---- 4. Train-Test Split ----
X = df.drop('fitness_status', axis=1)
y = df['fitness_status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# ---- Output ----
print("Processed Dataset Sample:")
print(df.head())

print("\nTrain Shape:", X_train.shape)
print("Test Shape:", X_test.shape)

```

Output:

```

... Processed Dataset Sample:
  student_id  height  weight  BMI  medical_condition \
0          1  1.791095 -4.125904e-01 -0.213982      0
1          2  0.000000  5.930988e-01 -1.394570      0
2          3 -1.399293 -1.418280e+00  0.671460      2
3          4 -0.335830  1.429170e-15  0.000000      0
4          5  1.791095 -4.125904e-01  0.000000      2

  fitness_status
0              0
1              0
2              0
3              1
4              0

Train Shape: (40, 5)
Test Shape: (10, 5)

```

Justification:

- Missing Value Imputation (Mean)
 - Maintains dataset size while preserving average distribution of health metrics.
- Categorical Encoding
 - Converts health conditions and fitness status into numerical form for predictive modeling.
- Standard Scaling

- Centers data (mean=0, std=1), essential for algorithms sensitive to feature magnitude (e.g., SVM, KNN).
- Train-Test Split
 - Separates data to evaluate model performance on unseen data, ensuring generalization.

Task 4: Real-Time Application: Ride-Sharing Data Cleaning

Scenario:

A ride-sharing company has trip data with errors.

Requirements:

- Standardize pickup and drop locations
- Fill missing fare amounts with median values
- Remove duplicate ride entries
- Normalize driver ratings
- Generate a summary of data improvements

Prompt:

Clean a ride-sharing dataset.

Instructions:

- Standardize pickup and drop locations to lowercase
- Fill missing fare values using median
- Remove duplicate records based on ride_id
- Normalize driver_rating using MinMax scaling
- Generate summary of cleaning steps and final dataset

Code:

```
# Task 4: Ride-Sharing Data Cleaning

import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

# ---- Sample Dataset ----
np.random.seed(2)

locations = ["Chennai", "Adyar", "T Nagar", "Velachery", "Tambaram"]
```

```

df = pd.DataFrame({
    "ride_id": np.random.randint(1, 40, 60),
    "pickup": np.random.choice(locations, 60),
    "drop": np.random.choice(locations, 60),
    "fare": np.random.choice([100, 150, 200, 250, np.nan], 60),
    "driver_rating": np.random.uniform(3.0, 5.0, 60)
})

# Introduce inconsistencies
df.loc[0, 'pickup'] = "chennai"
df.loc[1, 'drop'] = "adyar"

# ---- 1. Standardize Locations ----
df['pickup'] = df['pickup'].str.lower()
df['drop'] = df['drop'].str.lower()

# ---- 2. Handle Missing Fare ----
df['fare'] = df['fare'].fillna(df['fare'].median())

# ---- 3. Remove Duplicate Rides ----
before_count = len(df)
df = df.drop_duplicates(subset='ride_id')
after_count = len(df)

# ---- 4. Normalize Driver Ratings ----
scaler = MinMaxScaler()
df['driver_rating'] = scaler.fit_transform(df[['driver_rating']])

# ---- 5. Summary ----
print("CLEANED DATASET SAMPLE:")
print(df.head())

print("\nSUMMARY:")
print("Records Before Cleaning:", before_count)
print("Records After Removing Duplicates:", after_count)
print("Missing Values:\n", df.isnull().sum())

```

Output:

```
print("Missing Values:\n", df.isnull().sum())
```

```
... CLEANED DATASET SAMPLE:
   ride_id  pickup      drop  fare  driver_rating
0       16  chennai  chennai  200.0         0.976876
1        9   t nagar   adyar  100.0         0.787486
2       23  tambaram  chennai  100.0         0.423726
3       19   t nagar  velachery 250.0         0.341723
4       12  velachery   adyar  150.0         0.539948

SUMMARY:
Records Before Cleaning: 60
Records After Removing Duplicates: 28
Missing Values:
  ride_id      0
pickup      0
drop        0
fare        0
driver_rating 0
dtype: int64
```

Justification:

Location Standardization

- Ensures consistent grouping and analysis (e.g., “Chennai” vs “chennai”).

Missing Fare Imputation (Median)

- Median prevents distortion caused by unusually high/low fares.

Duplicate Removal

- Avoids repeated ride entries, ensuring accurate operational metrics.

Rating Normalization

- Scales ratings for fair comparison and compatibility with ML models.

Summary Generation

- Provides transparency and validates improvements in data quality.